

CSE 240A: Principles in Computer Architecture

ISA

Jishen Zhao (<https://cseweb.ucsd.edu/~jzhao/>)

[Adapted in part from Dean Tullsen, Mary Jane Irwin, Joe Devietti, Onur Mutlu, and others]

Review: Performance

- **Performance metrics:** Latency & throughput
- **Comparing performance:** Speedup
- **Averaging performance:**
 - Arithmetic mean
 - Harmonic mean
 - Geometric mean
- **Measuring CPU performance:**
 - CPI (IPC)
 - MIPS
 - FLOPS (later)
- **Performance laws**
 - Amdahl's law
 - Little's law
- **Benchmarks**
 - Micro-benchmarks / benchmarks


Review: Amdahl's Law

- Amdahl's Law:

How much does
an optimization improve performance?

$$\frac{1}{(1 - P) + \frac{P}{S}}$$

Amdahl's Law for the US Budget





Scrapping Dept of Transportation (\$81B) cuts budget by 2.6%

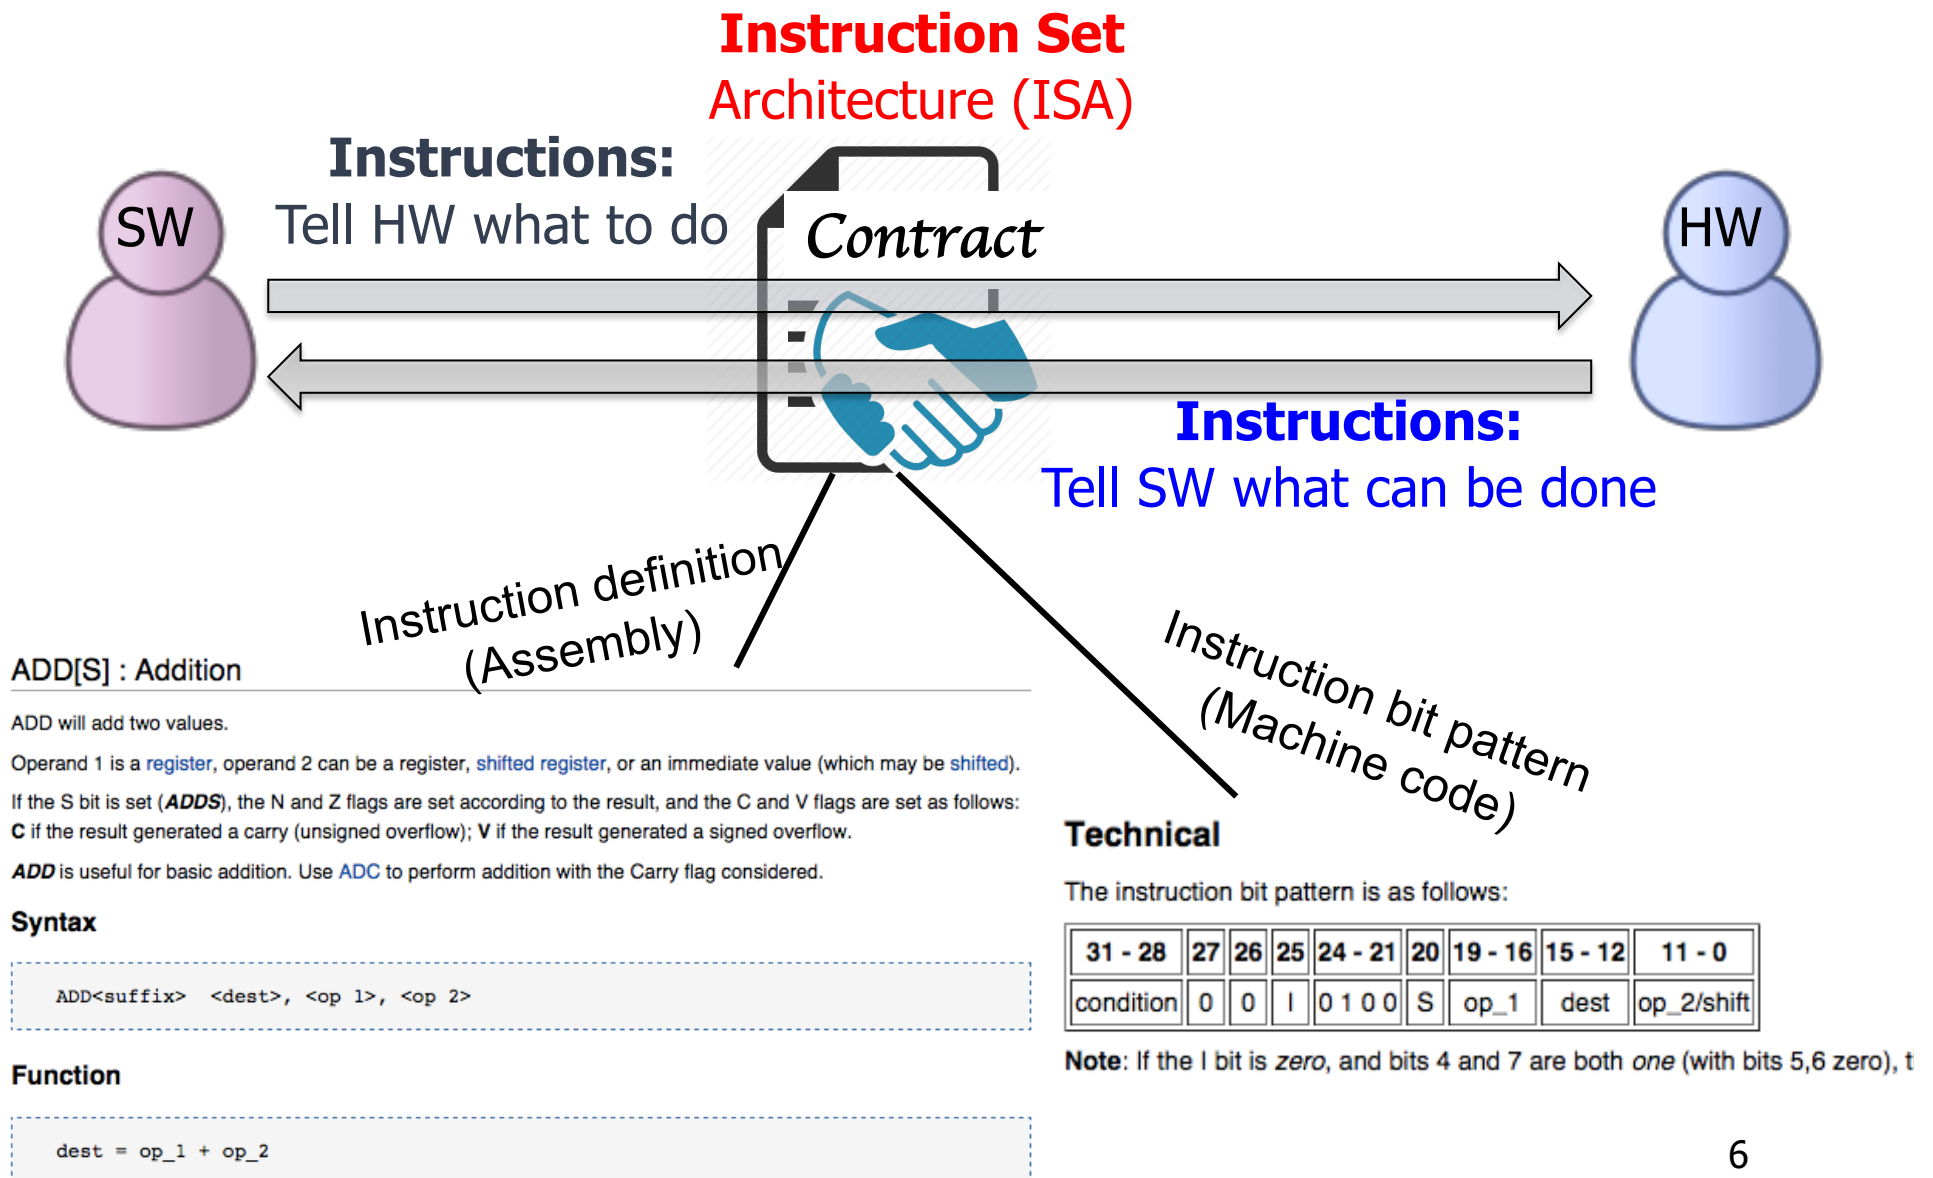
https://en.wikipedia.org/wiki/2015_United_States_federal_budget

Review: ISA

- What is ISA?
- Execution model:
 - Compilation
 - Assembly & machine language
- Instruction execution model
 - Registers, memory, PC
 - Instruction execution
- ISA design goals
 - Programmability
 - Performance/implementability
 - Compatibility



Review: What is ISA?




Review: what is/isn't defined in ISA?

- What's defined in ISA?
 - **Functional definition** of storage locations & operations
 - Operations: add, multiply, branch, load, store, etc
 - Data storage locations: registers, memory
 - **Precise description** of how to invoke operations & access data
- What's not in ISA? non-functional aspects
 - How operations are **implemented**
 - Which operations are **fast** and which are **slow** and when
 - Which operations take more **power** and which take less

Review: Instruction Execution Model


Instruction execution model
≠
Program execution model



- A computer executes **instructions**
 - **Fetches** next instruction from memory
 - **Decodes** it (figure out what it does)
 - **Reads** its **inputs** (registers & memory)
 - **Executes** it (adds, multiply, etc.)
 - **Write** its **outputs** (registers & memory)
 - **Next insn** (adjust the program counter)
- A computer is just a finite state machine
 - **Registers** (few of them, but fast)
 - **Memory** (lots of memory, but slower)
 - **Program counter** (next insn to execute)
 - Called "instruction pointer" in x86
- **Program is just "data in memory"**
 - Makes computers programmable ("universal")

Today: Aspects of ISA

- Instruction length
 - Next instruction: $PC + \text{length}$
- Instruction format
 - Instruction encoding
- Where does data live?
 - Addressing modes
- Control transfers
 - How to find the next instruction
 - Branch
 - Jump
- How to design high-performance ISA (if have time)
 - #1 make common case faster
 - #2 make fast case common



More on RISC vs. CISC

CISC and RISC adoptions

- The CISCiest: VAX (**V**irtual **A**ddress **eX**tension to PDP-11)
 - Variable length instructions: 1-321 bytes!!!
 - 14 registers + PC + stack-pointer + condition codes
 - Data sizes: 8, 16, 32, 64, 128 bit, decimal, string
 - Memory-memory instructions for all data sizes
 - Special insns: `crc`, `insque`, `polyf`, and a cast of hundreds
- x86: “Difficult to explain and impossible to love”
 - variable length insns: 1-15 bytes
- The RISCs: MIPS, PA-RISC, SPARC, PowerPC, Alpha, ARM
 - 32-bit instructions
 - 32 integer registers, 32 floating point registers
 - Load/store architectures with few addressing modes
 - Why so many basically similar ISAs? Everyone wanted their own

The RISC vs. CISC Design Tenets

- **RISC: Single-cycle execution**
 - CISC: many multicycle operations
- **RISC: Hardwired (simple) control**
 - CISC: **microcode** for multi-cycle operations
- **RISC: Load/store architecture**
 - CISC: register-memory and memory-memory
- **RISC: Few memory addressing modes**
 - CISC: many modes
- **RISC: Fixed-length instruction format**
 - CISC: many formats and lengths
- **RISC: Reliance on compiler optimizations**
 - CISC: hand assemble to get good performance
- **RISC: Many registers** (compilers can use them effectively)
 - CISC: few registers

Intel's x86 Trick: RISC Inside


- 1993: Intel wanted “out-of-order execution” in Pentium Pro
 - Hard to do with a coarse grain ISA like x86
- Solution? Translate x86 to RISC micro-ops (**μops**) in hardware
 - `push $eax`
becomes (we think, uops are proprietary)
`store $eax, -4($esp)`
`addi $esp, $esp, -4`
 - + Processor maintains **x86 ISA externally for compatibility**
 - + But executes **RISC μISA internally for implementability**
 - Given translator, x86 almost as easy to implement as RISC
 - Intel implemented “out-of-order” before any RISC company
 - “out-of-order” also helps x86 more (because ISA limits compiler)
 - Also used by other x86 implementations (AMD)
 - Different **μops** for different designs
 - **Not part of the ISA specification**, not publically disclosed

Potential Micro-op Scheme

- Most instructions are a **single** micro-op
 - Add, xor, compare, branch, etc.
 - Loads example: `mov -4(%rax), %ebx`
 - Stores example: `mov %ebx, -4(%rax)`
- Each memory access adds a micro-op
 - `"addl -4(%rax), %ebx"` is two micro-ops (load, add)
 - `"addl %ebx, -4(%rax)"` is three micro-ops (load, add, store)
- Function call (CALL) – 4 uops
 - Get program counter, store program counter to stack, adjust stack pointer, unconditional jump to function start
- Return from function (RET) – 3 uops
 - Adjust stack pointer, load return address from stack, jump register
- Again, just a basic idea, micro-ops are specific to each chip

Aspects of ISAs

Length and Format



- **Length**

- Fixed length
 - Most common is 32 bits
 - + Simple implementation (next PC often just PC+4)
 - Code density: 32 bits to increment a register by 1
- Variable length
 - + Code density
 - x86 averages 3 bytes (ranges from 1 to 15)
 - Complex fetch (where does next instruction begin?)
- Compromise: two lengths
 - E.g., MIPS16 or ARM's Thumb (16 bits)

- **Encoding**

- A few simple encodings simplify decoder
- Machine code (1s and 0s) <-> assembly

Example Instruction Encodings

- MIPS

- Fixed length
 - 32-bits, 3 formats, simple encoding
- add R1, R2, R3




- x86

- Variable length encoding (1 to 15 bytes)

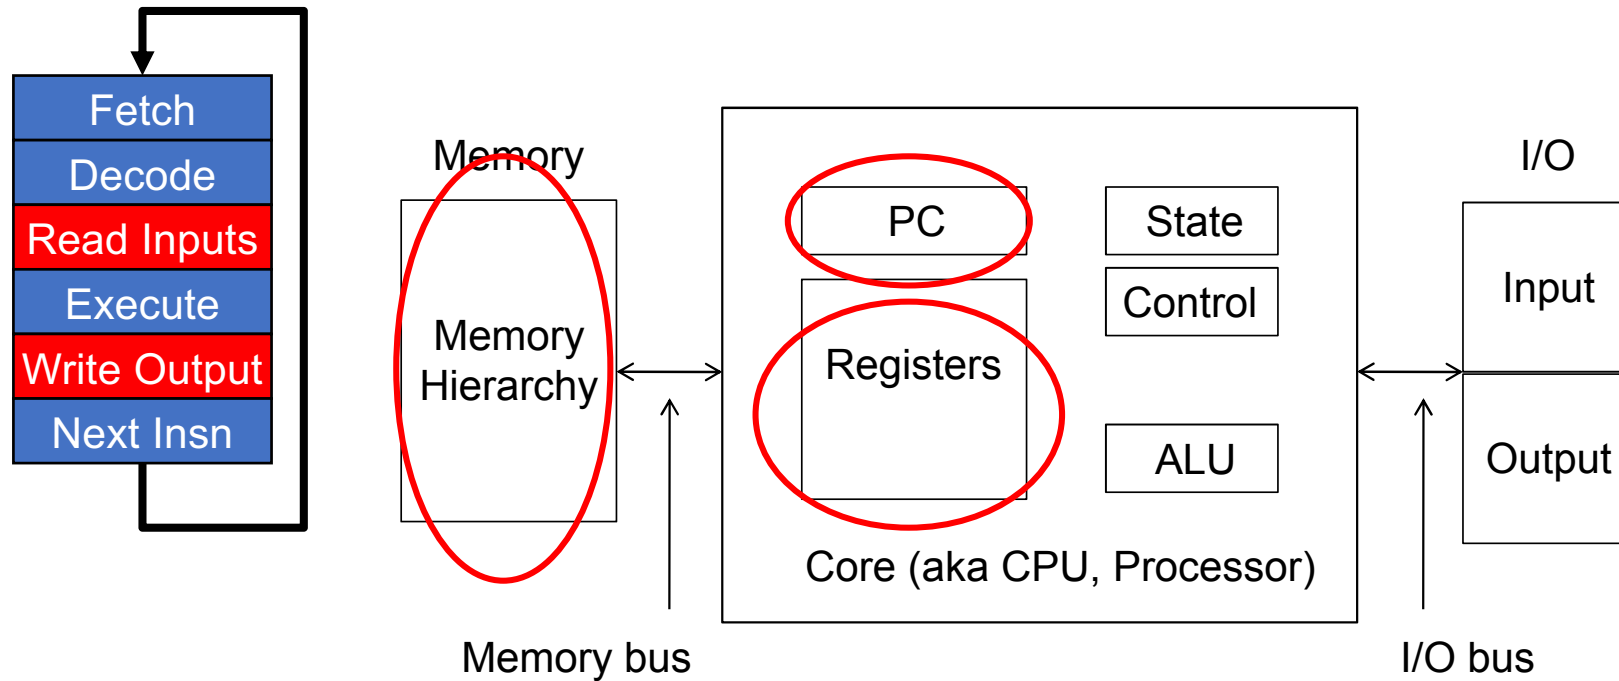


Operations and Datatypes




- Datatypes
 - Software: attribute of data
 - Hardware: attribute of operation, data is just 0/1's
- All processors support
 - Integer arithmetic/logic (8/16/32/64-bit)
 - IEEE754 floating-point arithmetic (32/64-bit)
- More recently, most processors support
 - "Packed-integer" insns, e.g., MMX
 - "Packed-floating point" insns, e.g., SSE/SSE2/AVX
 - For "data parallelism", more about this later
- Other, infrequently supported, data types
 - Decimal, other fixed-point arithmetic

Where Does Data Live?



Where Does Data Live?



- **Registers** (e.g., R0, R1, **F0**)

- "short term memory"
- Faster than memory, quite handy
- Named directly in instructions

ADD R1, R2, R3

- **Memory** (e.g., (R3), **#20(R5)**)

- "longer term memory"
- Accessed via "addressing modes"
 - Address to read or write calculated by instruction

ADD R1, R2, (R3)

- "Immediates" (e.g., #36, #7)

- Values spelled out as bits in instructions
- Input only

How Many Registers?

- Registers faster than memory, have as many as possible?
 - **No**
- One reason registers are faster: there are **fewer of them**
 - Small is fast (hardware truism)
- Another: they are **directly addressed** (no address calc)
 - More registers, means more bits per register in instruction
 - Thus, fewer registers per instruction or larger instructions
- **Not everything can be put in registers**
 - Structures, arrays, anything pointed-to
 - Although compilers are getting better at putting more things in
- More registers means **more saving/restoring**
 - Across function calls, traps, and context switches
- Trend toward more registers:
 - 8 (x86) → 16 (x86-64), 16 (ARM v7) → 32 (ARM v8)



```
0x0004
0x0003
0x0002
0x0001
```



0101	0110
1001	0101
0001	0111


Memory Addressing

- **Addressing mode:** way of specifying address
- Examples
 - **Displacement:** $\text{address} = [\text{R2} + \text{immed}]$, e.g., $\#20(\text{R2})$
 - **Index-base:** $\text{address} = [\text{R2} + \text{R3}]$
 - **Memory-indirect:** $\text{address} = [\text{mem}[\text{R2}]]$
 - **Auto-increment:** $\text{address} = [\text{R2}]$, $\text{R2} = \text{R2} + 1$
 - **Auto-indexing:** $\text{address} = [\text{R2} + \text{immed}]$, $\text{R2} = \text{R2} + \text{immed}$
 - **Scaled:** $\text{address} = [\text{R2} + \text{R3} * \text{immed1} + \text{immed2}]$
 - **PC-relative:** $\text{address} = [\text{PC} + \text{imm}]$

Addressing Modes Examples

- MIPS

I-type



- **Displacement:** R1+offset (16-bit)
- Why? Experiments on VAX (ISA with every mode) found:
 - 80% use small displacement (or displacement of zero)
 - Only 1% accesses use displacement of more than 16bits
- Other ISAs (SPARC, x86) have reg+reg mode, too
 - Impacts both implementation and insn count? (How?)
- x86 (MOV instructions)
 - **Absolute:** zero + offset (8/16/32-bit)
 - **Register indirect:** R1
 - **Displacement:** R1+offset (8/16/32-bit)
 - **Indexed:** R1+R2
 - **Scaled:** $R1 + (R2 * \text{Scale}) + \text{offset}(8/16/32\text{-bit})$ Scale = 1, 2, 4, 8

Example: x86 Addressing Modes

```
.LFE2
```

```
.comm array,400,32
```

```
.comm sum,4,4
```

```
.globl array_sum
```

```
array_sum:
```

```
movl $0, -4(%rbp)
```

Displacement



```
.L1:
```

```
movl -4(%rbp), %eax
```

```
movl array(,%eax,4), %edx
```

```
movl sum(%rip), %eax
```

```
addl %edx, %eax
```

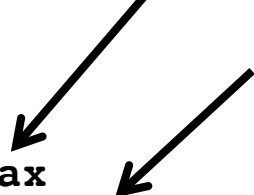
```
movl %eax, sum(%rip)
```

```
addl $1, -4(%rbp)
```

```
cmpl $99, -4(%rbp)
```

```
jle .L1
```

Scaled: address = array + (%eax * 4)
Used for sequential array access




PC-relative: offset of sum wrt %rip



Note: "mov" can be load, store, or reg-to-reg move


Control Transfers



- Default next-PC = PC + sizeof(current insn)
 - Branches and jumps can change that
- **Computing targets:** where to jump to
 - For all branches and jumps
 - **PC-relative:** e.g., bne R3, R6, L3
for branches and jumps with function
 - **Absolute:** e.g., J L3
for function calls
 - **Register indirect:** e.g., JR R5
for returns, switches & dynamic calls

```
L3:  
    addu   R7, R4, R3  
    lw     R7, (R7)  
    addu   R8, R5, R3  
    J      L3  
    bne    R3, R6, L3
```

Control Transfers



- **Testing conditions:** whether to jump or not
 - Implicit condition codes or “flags” (ARM, x86)

```
cmp R1,10    // sets  
              “negative” flag  
branch-neg target
```
 - Use registers & separate branch insns (MIPS)

```
set-less-than R2,R1,10  
branch-not-equal-zero  
R2,target
```

```
L3:  
    addu R7, R4, R3  
    lw  R7, (R7)  
    addu R8, R5, R3  
    J   L3  
    bne R3, R6, L3
```

ISAs Also Include Support For...

- Function calling conventions
 - Which registers are saved across calls, how parameters are passed
- Operating systems & memory protection
 - Privileged mode
 - System call (TRAP)
 - Exceptions & interrupts
 - Interacting with I/O devices
- Multiprocessor support
 - “Atomic” operations for synchronization
- Data-level parallelism
 - Pack many values into a wide register
 - Intel’s SSE2: four 32-bit float-point values into 128-bit register
 - Define parallel operations (four “adds” in one cycle)

ISA Code Examples

Code examples

```
int foo(int x, int y) {  
    return (x+10) * y;  
}
```

```
int max(int x, int y) {  
    if (x >= y) return x;  
    else return y;  
}
```

```
int array[100];  
int sum;  
void array_sum() {  
    for (int i=0; i<100;i++) {  
        sum += array[i];  
    }  
}
```

check out
<http://gcc.godbolt.org> to
examine these snippets

x86 and ARM, -O0 and -O3
x86: destination reg on the right
arm: destination reg on the left

How to design high-performance ISA?

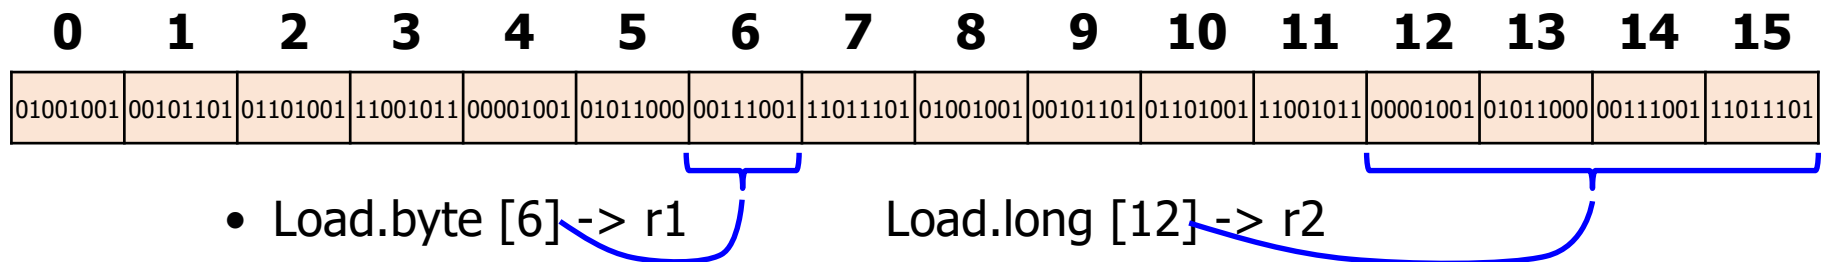
Performance Rule #1

make the common case fast

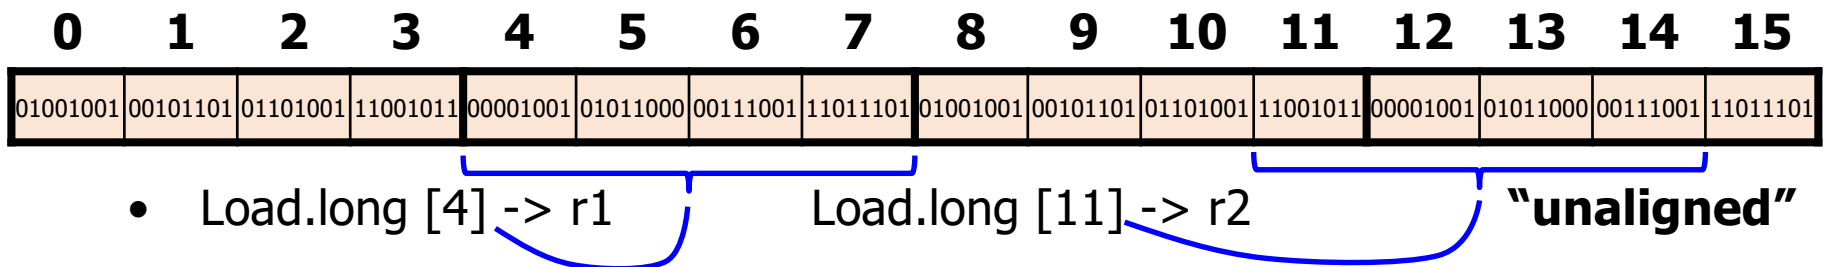
Access Granularity & Alignment

- **Byte addressability**

- An address points to a byte (8 bits) of data
- The ISA's minimum granularity to read or write memory
- ISAs also support wider load/stores
 - "Half" (2 bytes), "Longs" (4 bytes), "Quads" (8 bytes)



However, physical memory systems operate on **even larger chunks**



- **Access alignment:** if $\text{address} \% \text{size}$ is not 0, then it is "unaligned"
 - A single unaligned access may require multiple physical memory accesses

Handling Unaligned Accesses

- **Access alignment:** if $\text{address \% size} \neq 0$, then it is “unaligned”
 - A single unaligned access may require multiple physical memory accesses
- How to handle such unaligned accesses?
 1. Disallow (unaligned operations are considered illegal)
 - MIPS, ARMv5 and earlier took this route
 2. Support in hardware? (allow such operations)
 - x86, ARMv6+ allow regular loads/stores to be unaligned
 - Unaligned access still slower, adds significant hardware complexity
 3. Trap to software routine? (allow, but hardware traps to software)
 - Simpler hardware, but high penalty when unaligned
 4. In software (compiler can use regular instructions when possibly unaligned)
 - Load, shift, load, shift, and (slow, needs help from compiler)
 5. MIPS ISA support: unaligned access by compiler using two instructions
 - Faster than above, but still needs help from compiler

```
lwl @XXXX10; lwr @XXXX10
```

How big is this struct?


```
struct foo {  
    char c;  
    int i;  
}
```



Hint: avoid unaligned accesses

Another Addressing Issue: Endian-ness

- **Endian-ness**: arrangement of bytes in a multi-byte number
 - Big-endian: sensible order (e.g., MIPS, PowerPC, ARM)
 - The most significant byte (the "big end") of the data is placed at the byte with the lowest address
 - A 4-byte integer: "00000000 00000000 00000010 00000011" is 515_{10}
 - Little-endian: reverse order (e.g., x86)
 - A 4-byte integer: "00000011 00000010 00000000 00000000" is 515_{10}
 - The least significant byte (the "little end") of the data is placed at the byte with the lowest address
 - Why little/big endian?



Little-endian: Integer casts are free

starting address

Big-endian: Sign checks cheaper ← The sign bit is the most significant bit, and thus will be in the last byte in a little-endian format

Operand Model: Register or Memory?

- “Load/store” architectures
 - Memory access instructions (loads and stores) are distinct
 - Separate addition, subtraction, divide, etc. operations
 - Examples: MIPS, ARM, SPARC, PowerPC
- Alternative: mixed operand model (x86, VAX)
 - Operand can be from register **or** memory
 - x86 example: `addl 100, 4(%eax)`
 - 1. Loads from memory location `[4 + %eax]`
 - 2. Adds “100” to that value
 - 3. Stores to memory location `[4 + %eax]`
 - Would requires three instructions in MIPS, for example.

x86 Operand Model: Accumulators

```
.LFE2
```

```
.comm array,400,32
```

```
.comm sum,4,4
```

```
.globl array_sum
```

```
array_sum:
```

```
    movl $0, -4(%rbp)
```

```
.L1:
```

```
    movl -4(%rbp), %eax
```

```
    movl array(,%eax,4), %edx
```

```
    movl sum(%rip), %eax
```

```
    addl %edx, %eax
```

```
    movl %eax, sum(%rip)
```

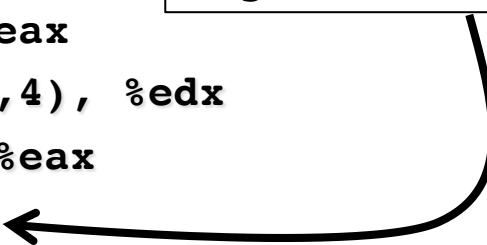
```
    addl $1, -4(%rbp)
```

```
    cmpl $99, -4(%rbp)
```

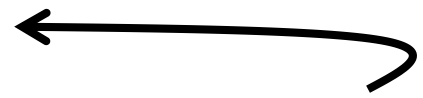
```
    jle .L1
```

- x86 uses explicit accumulators
 - Both register and memory
 - Distinguished by addressing mode

Register accumulator: $\%eax = \%eax + \%edx$



Memory accumulator:
 $\text{Memory}[\%rbp-4] = \text{Memory}[\%rbp-4] + 1$



How Much Memory? Address Size

- What does “64-bit” in a 64-bit ISA mean?
 - **Each program can address (i.e., use) 2^{64} bytes**
 - 64 is the size of **virtual address (VA)**
 - Alternative (wrong) definition: width of arithmetic operations
- Most critical, inescapable ISA design decision
 - Too small? Will limit the lifetime of ISA
 - May require nasty hacks to overcome (E.g., x86 segments)
- x86 evolution:
 - 4-bit (4004), 8-bit (8008), 16-bit (8086), 24-bit (80286),
 - 32-bit + protected memory (80386)
 - 64-bit (AMD’s Opteron & Intel’s Pentium4)
- All modern ISAs are at 64 bits

Performance Rule #2

make the fast case common

Winner for Desktops/Servers: CISC

- x86 was first mainstream 16-bit microprocessor by ~2 years
 - IBM put it into its PCs...
 - Rest is historical inertia, Moore's law, and "financial feedback"
 - x86 is most difficult ISA to implement and do it fast but...
 - Because Intel sells the most **non-embedded** processors...
 - It hires more and better engineers...
 - Which help it maintain competitive performance ...
 - **And given competitive performance, compatibility wins...**
 - So Intel sells the most **non-embedded** processors...
 - AMD has also added pressure, e.g., beat Intel to 64-bit x86
- Moore's Law has helped Intel in a big way
 - Most engineering problems can be solved with more transistors


Winner for Embedded: RISC

- ARM (Acorn RISC Machine → Advanced RISC Machine)
 - First ARM chip in mid-1980s (from Acorn Computer Ltd).
 - 6 billion units sold in 2010
 - Low-power and **embedded/mobile** devices (e.g., phones)
 - Significance of embedded? ISA compatibility less powerful force
- 64-bit RISC ISA
 - 32 registers, PC is one of them
 - Rich addressing modes, e.g., auto increment
 - Condition codes, each instruction can be conditional
- ARM does not sell chips; it licenses its ISA & core designs
- ARM chips from many vendors
 - Apple, Qualcomm, Freescale (né Motorola), Texas Instruments, STMicroelectronics, Samsung, Sharp, Philips, etc.

Redux: Are ISAs Important?

- Does “quality” of ISA actually matter?
 - Not for performance (mostly)
 - Mostly comes as a design complexity issue
 - Insn/program: everything is compiled, compilers are good
 - Cycles/insn and seconds/cycle: μ ISA, many other tricks
 - What about power efficiency? Maybe
 - ARMs are most power efficient today...
 - ...but Intel is moving x86 that way (e.g., Atom)
- Does “nastiness” of ISA matter?
 - Mostly no, only compiler writers and hardware designers see it
- Comparison is confounded by, e.g., transistor technology
- Even compatibility is not what it used to be
 - cloud services, virtual ISAs, interpreted languages

Instruction Set Architecture (ISA)



- What is an ISA?
 - A functional contract
- All ISAs similar in high-level ways
 - But many design choices in details
 - Two “philosophies”: CISC/RISC
 - Difference is blurring
- Good ISA...
 - Enables high-performance
 - At least doesn’t get in the way
- Compatibility is a powerful force
 - Tricks: binary translation, μ ISAs